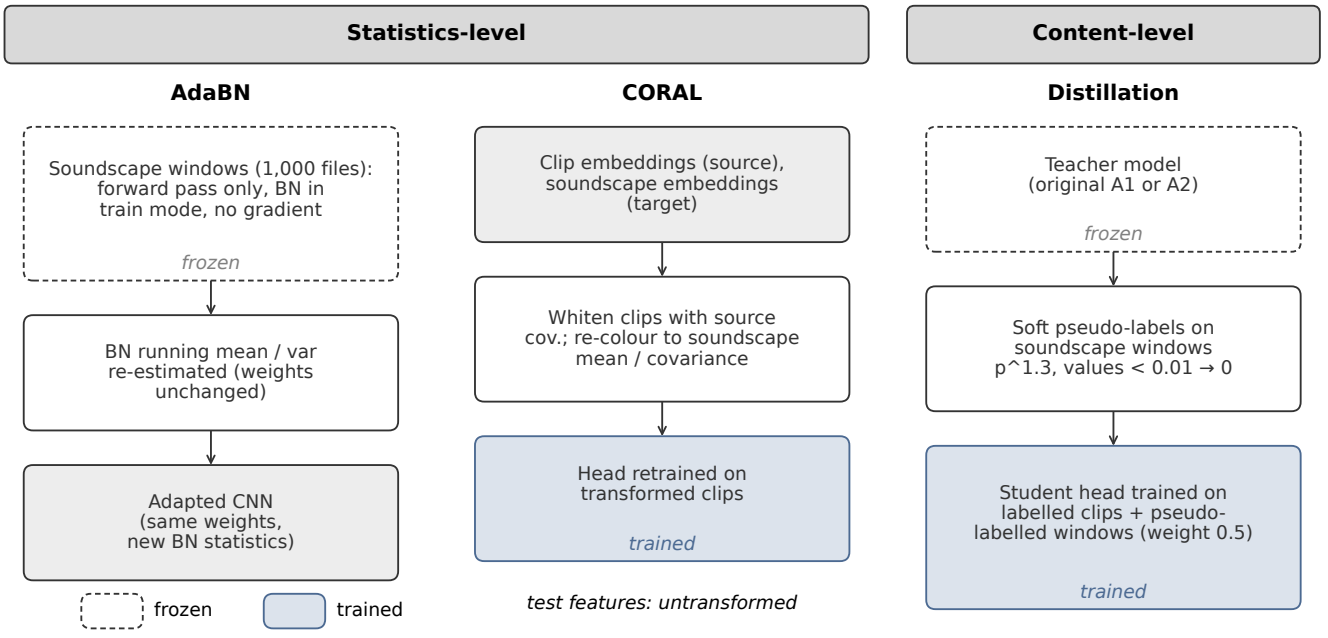
Statistics-level
Content-level
AdaBN
CORAL
Distillation
Soundscape windows (1,000 files): forward pass only, BN in train mode, no gradient
frozen
BN running mean / var re-estimated (weights unchanged)
Adapted CNN (same weights, new BN statistics)
frozen
trained
Clip embeddings (source), soundscape embeddings (target)
Whiten clips with source cov.; re-colour to soundscape mean / covariance
Head retrained on transformed clips
trained
test features: untransformed
Teacher model (original A1 or A2)
frozen
Soft pseudo-labels on soundscape windows p^1.3, values < 0.01 → 0
Student head trained on labelled clips + pseudo-labelled windows (weight 0.5)
trained